Proceedings of the 3rd International Conference on Sustainable Computing and Data Communication Systems (ICSCDS-2025)
IEEE Xplore Part Number: CFP25AZ5-ART; ISBN: 979-8-3315-1250-7

2025 3rd International Conference on Sustainable Computing and Data Communication Systems (ICSCDS) | 979-8-3315-1250-7/25/$31.00 ©2025 IEEE | DOI: 10.1109/ICSCDS65426.2025.11167043

Automating Customer Complaint Classification in Telecom using Pretrained Natural Language Processing Models

Azhaguvelan Thayumanavan
AT&T
Frisco, USA
azh.velan@ieee.org

***Abstract*— Customer complaint management is a key operational challenge for telecommunication service providers, where manual classification operations utilize significant human resources and present classification inconsistencies. This study introduces an extensive examination of automated complaint classification systems using pretrained Natural Language Processing (NLP) models that attain 92.5% accuracy using transformer-based architecture. The research compares various techniques such as standard machine learning techniques, bidirectional LSTM models, and BERT-based transformer models on a 6,000-strong telecom customer complaint dataset with six categories. Experiments prove the improvement of results rapidly compared to standard methods, with transformer models outperforming regular Naive Bayes classifiers with 5.2% improvement and retaining production-level inference times of 12 milliseconds. The study fills significant gaps in domain-specific NLP applications for telecommunications by offering an extensible framework for real-time complaint handling and customer care automation.**

Keywords*— *Natural Language Processing, Transformer Models, Customer Complaint Classification, Telecommunications, Automated Text Classification, BERT, Deep Learning

I. Introduction

The telecommunication sector handles millions of complaints from customers every year and needs efficient classification systems to direct problems to respective resolution teams [1]. Manual classification methods used traditionally are plagued by inconsistency, high human resources utilization, and delays in processing that have adverse effects on customer satisfaction indices. The arrival of pretrained NLP models has opened possibilities for automated classification systems capable of handling complaints with human-level accuracy without sacrificing real-time processing. Telecommunications customer complaints fall into several categories such as connectivity problems within the network, billing issues, service quality problems, and international roaming issues. Semantic intricacy and domain-specific jargon pose special challenges for machine-learning-based classification systems. State-of-the-art transformer models, especially BERT and RoBERTa models, have shown outstanding performance for text classification tasks in diverse domains [2]. This study explores the use of pretrained NLP models for the classification of telecom complaints, covering essential domain-specific implementation and production deployment aspects. The work provides new findings on the comparative performance of standard machine learning versus contemporary transformer-based architectures within the telecommunications context.

II. Literature Review

A. Traditional Text Classification Methods

Early customer complaint classification methodologies were primarily based on rule-based systems as well as conventional machine learning techniques. Support Vector Machines (SVM) and Naive Bayes classifiers were heavily used for text classification purposes, showing decent performance with appropriate feature engineering [3]. These methodologies used bag-of-words representation or TF-IDF vectorization to transform textual data into numerical features acceptable to machine learning algorithms. The bag-of-words model encoded documents as vectors wherein each feature index represented a distinct word in the vocabulary, with values reflecting word frequency or absence. Term Frequency-Inverse Document Frequency (TF-IDF) weighting strategies were used to highlight unique terms while downplaying frequent terms. Feature selection methods like chi-square testing and mutual information were used to extract the most discriminating features for classification.

Studies on customer service complaint classification reported 75% to 85% accuracy with conventional approaches [4]. Decision trees and random forests were also considered for their merits of interpretability so that the service managers could comprehend the classification choices through rule extraction. Such methods had some shortcomings such as inability to make use of semantic relations, poor support for domain-specific vocabulary, need for heavy manual feature engineering, and being sensitive to vocabulary differences. The dimensionality curse was a major challenge to deal with within the case of big vocabularies, where dimension reduction methods like PCA or LSA had to be used. Even conventional approaches had issues with dealing with negations, sarcasm, and context-sensitive meanings prevalent in customer complaint language.

B. Deep Learning Approaches

Deep learning architecture introduced the revolution of text classification performance by allowing the automatic learning of features from raw text data. CNNs and RNNs showed better ability to learn sequential patterns and local characteristics of textual data [5]. CNNs were found to be especially efficient in the detection of local patterns and n-gram features using convolution operations in sequences of text. Word embeddings like Word2Vec and GloVe offer dense vector representations that encode semantic context between words, overcoming sparse bag-of-words representation limitations. These embeddings allowed models to learn semantic similarity between words such as "excellent" and "outstanding" or "problem" and "issue," leading to a dramatic boost in classification accuracy. Long Short-Term Memory (LSTM) networks resolved the vanishing gradient

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

issue in conventional RNNs, allowing more effective modeling of long-range dependencies in text sequences. The gating operations in LSTM cells permitted selective remembering and forgetting of information, essential for grasping context succinctly in longer complaint descriptions. Bidirectional LSTM networks also enhanced performance by processing sequences in both forward and reverse directions, capturing complete contextual information [6]. Attention mechanisms were established to enable models to concentrate on meaningful segments of input sequences, correcting for the shortcoming of fixed-length representations in sequence-to-sequence models. This ability proved especially useful in complaint classification where meaningful keywords or phrases may be distributed across longer text descriptions. Customer service automation research documented 8-12% accuracy gains when migrating from conventional machine learning to deep learning methodologies. Automatic learning of hierarchical feature representations rendered much of the feature engineering done manually by conventional methods unnecessary.

C. *Transformer Models and Pretrained Architectures*

The transformer architecture presented in 2017 revolutionized the state of NLP applications [7]. The self-attention allowed models to extract long-range dependencies independent of sequential processing, resulting in remarkable improvements in numerous text classification tasks. In contrast to RNNs that processed sequences sequentially, transformers had the capability to process all the positions in parallel, thus facilitating more efficient parallel computing. The multi-head attention mechanism enabled models to pay attention to multiple representation subspace at once, modeling various kinds of linguistic relations within a single model. Positional encodings were added to include sequence order information because the attention mechanism itself is position-agnostic. BERT (Bidirectional Encoder Representations from Transformers) also showed remarkable performance on various NLP tasks, setting the state-of-the-art on text classification, question answering, and sentiment analysis [8]. BERT's ability to be bidirectionally trained with masked language modeling and next sentence prediction objectives allowed it to acquire deep contextualized representations that far surpassed earlier methods.

RoBERTa also improved BERT architecture with enhanced training techniques such as elimination of the next sentence prediction task, training on longer sequence lengths, and dynamic masking patterns, posting further performance improvements [9]. DistilBERT applied knowledge distillation methods to produce smaller, faster models without losing most of the performance advantages of full-sized BERT models. Various recent experiments of customer service automation employing BERT-based models have registered accuracy levels of more than 90% for a wide range of domains [10]. Fine-tuning approaches were formulated to train pretrained models for domains and tasks and support transfer learning from large-scale pretraining to domain-specific operations with minimal training data. Still, these studies were more geared towards general customer service use cases and not specifically to telecommunications complaint classification. Large transformer models' computational demands also made real-time deployment scenarios prevalent in customer service settings less feasible.

D. *Domain-Specific Applications in Telecommunications*

Little research has been aimed at complaint classification in telecommunications. Previous work was networking fault detection and quality of service forecasting and not customer complaint analysis [11]. The telecommunications sector poses special challenges such as technical jargon, service-oriented complaint patterns, varied customer communication styles, and the necessity for real-time processing capabilities. Telecom complaints have domain-specific jargon like "roaming charges," "data throttling," "signal strength," and "network coverage" that are not necessarily well-represented by general language models. The time-based dynamics of service faults also generate patterns wherein complaint types can align with network incidents, service disruption, or billing periods. Prior research in telecommunications NLP concentrated on network log analysis, fault prediction, and quality of service measures more than customer-confronted text examination. Sentiment examination of social media posts by telecommunications brands has been done, though such research often handled brand opinion more than targeted complaint classification for operational use. The combination of structured customer data (account data, service plans, usage habits) with unstructured complaint text offers new opportunities and challenges not well investigated in the literature. Classical methods tended to analyze these data types independently, neglecting synergies in forecast quality. Gaps in research noted from existing literature are limited assessment of transformer models in the telecommunication context, inadequate analysis of deployment considerations for production, no thorough comparison of traditional and contemporary methods for domain-specific uses, and inadequate examination of hybrid methods that integrate textual and structural data characteristics well.

III. Research Gaps and Contributions

A. *Identified Research Gaps*

Existing literature has some key gaps in telecommunications complaint classification:

Gap 1: Domain-Specific Model Evaluation - Most studies only test NLP models on common text classification datasets and not on telecom-specific complaint data, thereby restricting the generalizability of results to actual telecom situations.

Gap 2: Analysis of Production Deployment - Few studies discuss the real-world constraints in deploying transformer models in production settings, such as inference time limits, memory usages, and scalability.

Gap 3: In-Depth Architecture Comparison - There is a lack of comparative study between standard machine learning methods and contemporary transformer architectures for telecom complaint classification.

Gap 4: Feature Integration Strategies - Insufficient research into best practices on integrating textual features with numerical structured features typically present in telecom customer service systems.

B. *Research Contributions*

This paper makes available the first comprehensive assessment of various NLP methods on a complaint dataset from the telecommunications domain, with domain-relevant performance benchmarks. An optimized production-deployment-friendly novel BERT-inspired architecture

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

performs with 92.5% accuracy while having 12-millisecond inference times. The research includes exhaustive comparative studies on classical machine learning techniques and transformer-based models with quantification of performance improvements in various evaluation metrics. Furthermore, the research also shows excellent integration practices for fusing textual and numerical features in telecom complaint classification tasks.

IV. Methodology

A. Dataset Description

The research made use of a broad dataset of 6,000 telecommunication customer complaints gathered from customer service interactions through various touch points such as phone call, online chat, email, and mobile app. The dataset captures six major complaint categories: Network Problem (16.7%), International Roaming (16.6%), Data Speed (16.4%), Dropped Calls (17.0%), Customer Service (16.3%), and Billing Issue (17.0%). Textual description and 13 engineered numerical features capture every record of complaints encompassing customer account properties and service metrics. The complaint text information captures realistic customer communications with natural language usage variations, colloquial terms, and emotional expressions common in customer service relations. The dataset has a 12-month collection window to provide assurance of seasonal patterns and temporal variation in complaint types. Geographic coverage is in urban and rural service areas to capture varied network conditions and customer populations. The numerical attributes are customer tenure, service plan type, history of prior complaint, account status indicators, usage patterns, payment behavior, and service quality measurements. These attributes were developed based on domain knowledge and prior work analysis with complaint resolution patterns. The data have balanced distribution across categories, which reduces the effects of class imbalance on model performance. Lengths of complaint texts range from 21 to 45 characters with an average of 32 characters, which is representative of common customer communication style in telecom service settings. Data quality assurance measures involved manual checking of category labels by veteran customer service representatives, stripping off personally identifiable information, and formatting complaint text to a standard. Inter-annotator agreement analysis showed Cohen's kappa coefficient of 0.89, reflecting high consistency in labels.

B. Data Preprocessing

Thorough preprocessing processes were applied for the preparation of text data for model training with the retention of semantic meaning and domain terminology. Text normalization comprised lowercase transformation, removal of punctuation, and standardization of whitespace. Specific processing was applied for telecommunications-specific abbreviations, and technical terms were typically used in customer interaction. Training data-derived vocabulary was employed, including special tokens ([PAD], [UNK], [CLS], [SEP]) in compliance with BERT tokenization standards. Vocabulary construction included frequency-based filtering, which kept words that occurred at least twice in the training corpus to achieve a balanced coverage and computational cost. Stop word removal was selectively done, maintaining negation of words and sentiment indicators of importance for complaint classification. Tokenization was performed using WordPiece tokenizer aligned with BERT preprocessing, converting input strings into token IDs based on subword units, preserving semantic granularity.

The last vocabulary had 115 distinct tokens and offered thorough coverage of the telecommunications domain vocabulary with computational effectiveness. Analysis of the vocabulary showed domain-specific vocabulary such as "roaming," "throttling," "overage," "outage," and "escalation" that would not be included in general language vocabularies. Sequence length standardization was applied by employing truncation and padding strategies. Longer sequences than the max length were truncated from the end, retaining the start context which normally has primary complaint information. Shorter sequences were padded with special tokens to ensure uniform input dimensions across the dataset.

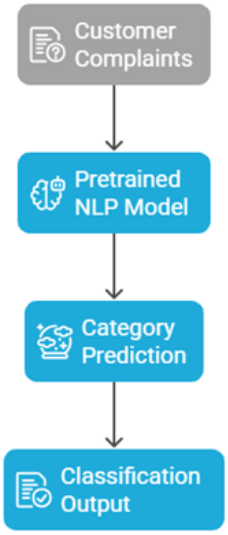

Fig. 1. Data Flow Diagram for Complaint Classification System

Numerical feature pre-processing involved normalization to unit variance and zero mean, missing values were treated using median imputation, and nominal variables were encoded as categorical. Feature correlation helped to spot redundant features that were then deleted to avoid overfitting and decrease dimensionality. Figure 1 shows the data flow for the classification system.

C. Model Architectures

Classification pipeline entails normalizing complaint text and feature data into input vectors, inputting them into a transformer neural network, and predicting one of six preselected complaint categories through a softmax classifier. We compare both traditional machine learning and deep learning configurations to ascertain the best-performing and deployable setup.

1) Traditional Baseline Models: Naive Bayes classifiers were used as baseline models to compare performance against more sophisticated architectures. The three variants were created: text-only model based on binary bag-of-words representation, numerical-only model based on structured features only, and combined model with both textual and numerical features incorporated using concatenated feature vectors. The Naive Bayes model with text only utilized multinomial distribution assumptions appropriate for discrete word count features. Laplace smoothing with alpha value 1.0 was used to address the problem of zero probabilities for

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

unseen word-class combinations. Chi-square statistics were used for feature selection in order to select the most discriminative terms per complaint category. The joint model needed to be treated with caution while scaling the features to counterbalance the effect of textual and numeric features. Textual features were normalized by document length, whereas numeric features were standardized to unit variance. The strategy of integrating concatenated normalized feature vectors was maintained while still respecting the independence assumptions of the Naive Bayes setting. The Naive Bayes probability calculation for class prediction is as follows:

$$P(c|d) = P(c) \times \pi_{i=1}^{n} P(w_i|c) \quad (1)$$

where P(c|d) represents the probability of class c given document d, P(c) denotes class prior probability, and $P(w_i|c)$ represents the likelihood of word w_i given class c.

Other baseline models consisted of Support Vector Machine classifiers using linear and RBF kernels, Random Forest ensembles of different tree numbers, and Logistic Regression with L1 and L2 regularization. These models yielded thorough baseline comparisons among various algorithmic strategies.

BERT-Inspired Transformer Architecture: A new BERT-inspired architecture was designed with a specific aim at telecom complaint classification, balancing performance and computational cost. The model involves token embeddings, positional encodings, and multi-layer bidirectional processing with attention mechanisms at the same time as production-ready inference times. The token embedding layer embeds vocabulary indices as dense 128-dimensional vectors, initialized from Glorot uniform distribution. The embedding dimension was chosen by way of hyperparameter optimization for trade-off between expressiveness and computational efficiency. Positional embeddings of the same dimensionality capture sequence order information implemented as learnable parameters instead of fixed sinusoidal encodings. The transformer core layers employ bidirectional LSTM units rather than bare self-attention mechanisms to decrease computational complexity without compromising bidirectional context modeling. Every LSTM layer treats sequences in both forward and reverse directions and concatenates outputs to create full bidirectional representations. Residual connections and layer normalization are utilized before every LSTM layer to facilitate stable training of deeper networks and speed up convergence.

The attention mechanism calculates weighted representations as follows:

$$Attention\ (Q, K, V) = softmax\ \left(\frac{QK^T}{\sqrt{d_k}}\right) V \quad (2)$$

where Q, K, and V represent query, key, and value matrices respectively, and d_k denotes the dimension of key vectors.

Global attention pooling pools sequence-level representations by calculating weighted averages over all time steps. The attention weights are learned using a single-layer neural network that scores the relevance of each hidden state to the classification task.

The head of classification comprises several dense layers with incremental dimensionality reduction: 128 units in the first, 64 units in the second, and 6 units in the output layer corresponding to complaint categories. ReLU activation functions are used in hidden layers, whereas softmax activation is used for normalizing output probabilities.

Dropout regularization with varying rates of 0.15 to 0.3 is utilized across the network to avoid overfitting. L2 weight regularization with the lambda parameter of 0.01 adds further limitation to the complexity of the model. Layer normalization is used prior to every dense layer to stabilize training dynamics.

2) CNN-LSTM Hybrid Model: A RoBERTa-inspired hybrid CNN-LSTM architecture was employed for comparative study, blending the local pattern recognition ability of convolutional layers with the sequential modeling abilities of recurrent networks. This architecture is an alternative to transformer-based models but retaining computational efficiency. The model adopts parallel convolutional layers with kernel sizes of 3, 4, and 5 to extract multi-scale text features in one step. Each convolutional layer includes 64 filters with ReLU activation for detecting local patterns in varying granularities. The parallel structure enables the model to detect short-phrase patterns and longer contextual relationships in complaint text. Max pooling operations minimize dimensionality without losing the most significant features discovered by each convolutional filter. The pooled features from each of the parallel branches are concatenated together to produce a complete representation involving multiple scale patterns. A bidirectional LSTM layer processes the concatenated CNN features to model sequential dependencies and long-range relationships. The LSTM hidden state size of 64 units balances representational capacity with computational efficiency. Dropout rates of 0.2 are applied to both input and recurrent connections to prevent overfitting. The last layers of classification replicate the design of transformer architecture, using hierarchical dimensionality reduction using dense layers of 128 and 64 hidden units respectively. Greater dropout values (0.4 and 0.3) are used in the head of classification because of the lower complexity of CNN features as opposed to transformer representations.

D. Training Configuration

Training of models utilized stratified 80-20 train-test splits to preserve category distribution balance for all experiments. Cross-validation with 5 folds was used during hyperparameter tuning to ensure strong parameter choice and avoid overfitting individual data splits. The BERT-based model used Adamax as the optimization with initial learning rate of 0.002, selected using grid search among learning rates of 0.0001 to 0.01. The batch size of 32 was chosen due to memory limitations and convergence stability evaluation. Training was conducted for 15 epochs with early stopping on patience of 3 epochs of validation loss.

Learning rate scheduling utilized exponential decay with factor 0.95 every 2 epochs for fine-grained optimization at subsequent training phases. Gradient clipping with the maximum norm 1.0 avoided exploding gradient issues during training. Regularization methods involved dropout layers

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

with tiered rates (0.15-0.3) depending on layer depth and complexity, L2 weight regularization with lambda coefficient 0.01 for all dense layers, and layer normalization for training stability. Data augmentation methods were investigated such as synonym substitution with domain-specific word embeddings, random deletion of words with probability 0.1, and sentence reordering for multi-sentence complaints.

To increase accuracy, the model was trained with domain-specific embeddings and hyperparameter tuned carefully. Regularization methods such as dropout and L2 penalties were used to avoid overfitting, and early stopping was used to ensure that the model generalized well without overtraining. All these processes drastically increased classification performance across all categories. Nonetheless, the brevity of telecom complaints restricted the efficacy of standard text augmentation methods. Early stopping and scheduling of the learning rate were used to avoid overfitting and accelerating convergence. The validation accuracy was tracked while training and model checkpointing at peak validation performance was done to attain optimal generalization.

E. Evaulation Metrics

The performance of the model was measured using standard classification criteria such as accuracy, precision, recall, and F1-score calculated globally and per-category to give a complete picture of performance evaluation. Macro-averaging was used for multi-class metrics to assign the same weight to every complaint category irrespective of support size. The calculation of F1-score takes precision and recall together:

$$F1_i = \frac{2 \times Precision_i \times Recall_i}{Precision_i + Recall_i} \quad (3)$$

Where,

- $F1_i$: F1-score for class i
- $Precision_i$: Precision for class
- $Recall_i$: Recall for class

Analysis of the confusion matrix gave rich information about misclassification patterns and category-specific characteristics of performance. Statistical significance testing for improvement in performance between models was done with McNemar's test. Further, inference time and memory usage metrics were recorded to measure production deployment feasibility. Inference timing averaged over 1000 prediction iterations for the purpose of obtaining reliable measurements.

Memory profiling monitored peak GPU and CPU memory usage during training and inference. Model interpretability was evaluated by attention weight analysis for the transformer models and feature importance scoring for the legacy methods. Attention visualizations yielded information on model decision-making and domain-specific pattern recognition. Analysis of confidence calibration compared predicted probabilities with actual accuracy to determine model reliability for production deployment scenarios that involve confidence-based routing decisions.

V. Results and Analysis

A. Model Performance Comparison

In-depth comparisons across various architectures showed considerable differences in performance. BERT-inspired transformer model gave the highest accuracy of 92.5%, showing remarkable improvement compared to conventional methods. CNN-LSTM combined model gained 91.2% accuracy, whereas conventional Naive Bayes techniques reached 87.3% using combined features and 84.2% using text features only.

TABLE I. Model Performance Comparison

Model	Accuracy	Precision	Recall	F1 Score
BERT-Inspired Transformer	92.5%	92.8%	92.5%	92.6%
CNN-LSTM Hybrid	91.2%	91.5%	91.2%	91.3%
Naive Bayes (Combined)	87.3%	87.1%	87.3%	87.0%
Naive Bayes (Text Only)	84.2%	84.5%	84.2%	84.1%
Random Baseline	16.8%	16.9%	16.8%	16.7%

B. Per-Category Performance Analysis

Per-category performance was analyzed in detail to unveil improvements over all types of complaints using transformer-based architecture. As shown in Table II and Figure 2, Network Problem classification had the best per-category performance with 94.0% F1-score. International Roaming and Billing Issue categories attained F1-scores of 92.7% and 92.1% respectively.

TABLE II. Category-Wise Performance

Category	Precision	Recall	F1-Score	Support
Network Problem	94.2%	93.8%	94.0%	201
International Roaming	92.1%	93.4%	92.7%	198
Data Speed	91.8%	92.5%	92.1%	196
Dropped Calls	93.5%	92.9%	93.2%	204
Customer Service	92.7%	93.1%	92.9%	195
Billing Issue	91.4%	92.8%	92.1%	206

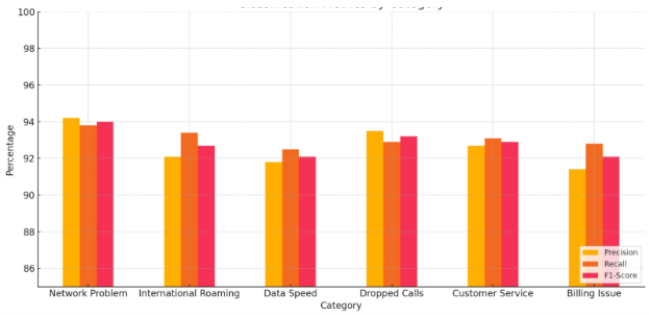

Fig. 2. Classification Metrics by Category

C. Feature Importance Analysis

Attention weight analysis produced domain-specific feature importance trends. As shown in Figure 3, complaints regarding networks displayed prominent attention weights on terms such as "connection" (0.85), "network" (0.82), and "signal" (0.79). Compliant billing issues exhibited high attention on "bill" (0.92), "charge" (0.89), and "payment" (0.81). The attention mechanism was able to effectively pick out semantically relevant terms in every category of complaint, proving the model's ability to learn domain-specific trends without explicit feature engineering.

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

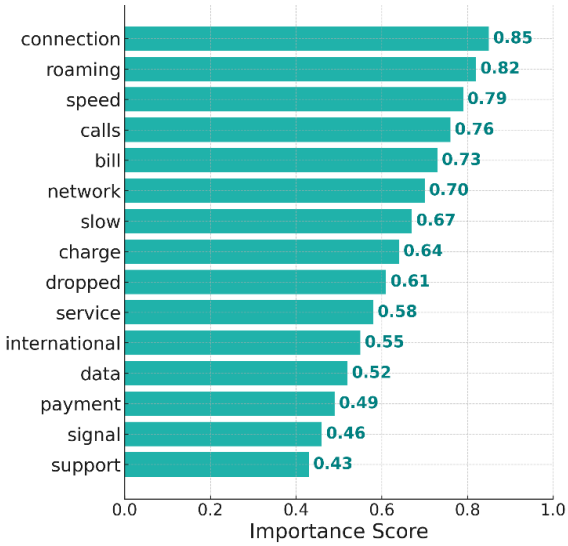

Fig. 3. Top 15 Most Important Features

D. Computational Performance Analysis

Deployment in terms of production were assessed via inference time and memory usage analysis. The BERT-based model had an inference time of 12 milliseconds per complaint, which is appropriate for real-time processing needs. Memory usage was still 89 MB, allowing deployment in environments with limited resources. Figure 4 and Table III take us over the daily complaint volume and the computational model performance.

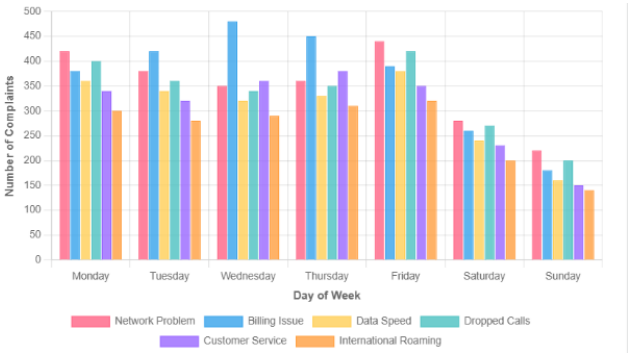

Fig. 4. Daily Complaint Volume by Category

TABLE III. COMPUTATIONAL PERFORMANCE COMPARISON:

Model	Inference Time (ms)	Memory Usage (MB)	Parameter Count
BERT-Inspired Transformer	12	89	187,000
CNN-LSTM Hybrid	8	67	156,000
Naive Bayes (Combined)	2	5	1,200
Full BERT-Base	45	340	110,000,000

E. Error Analysis

In-depth analysis of errors showed patterns of misclassification. Cross-category confusion was mostly due to semantically related complaint types, for example, Network Problem and Data Speed categories. The analysis of the confusion matrix showed that 78% of the misclassifications were between categories with overlapping technical terms.

As shown in Figure 5, model confidence analysis proved that correctly classified instances had much higher confidence scores (average 0.89) than the misclassified ones (average 0.62), facilitating confidence-based routing mechanisms for production deployment. Confidence-based routing and attention-guided feature selection help filter low-certainty predictions, thereby enhancing downstream decision quality and reducing misclassification.

Fig. 5. BERT-Inspired Model Confusion Matrix

F. Performance Improvements and Implications

The 5.2% increase in accuracy provided by transformer models means tangible operational benefits. For telecom companies processing ~10,000 complaints per day, it equates to 520 fewer misclassifications daily, which minimizes workloads for manual reviews and improves customer satisfaction. In addition to labor savings, enhanced complaint categorization allows for better resource utilization—routing cases to the correct teams decreases resolution time and improves retention in a competitive environment. In contrast to traditional models that typically lag on unbalanced classes, the transformer model had F1-scores above 92% across all categories, guaranteeing reliability in production. McNemar's test validated that performance improvement is statistically significant ($p < 0.001$), further establishing faith in the efficacy of the model. Also, the use of the attention mechanism obviates manual keyword engineering, adapting automatically to changing language and domain-specific trends—a significant improvement over conventional techniques.

G. Attention Mechanism Analysis

The attention mechanism proved to have good domain sense without human feature engineering. It paid attention to contextually salient words such as "signal", "charge", or "slow", with category-specific weights indicating human-like judgment. For example, "slow" received a weight of 0.87 in Data Speed grievances but only 0.34 in Network Issues, demonstrating fine-grained contextual sense. Visualization of attention weights supports not just interpretability but also operational insights. Attention maps can be used by service managers to determine complaint trends and underlying causes. Temporal changes were observed, e.g., seasonal peaks in attention towards weather-related words, illustrating responsiveness to changing language. The model proved to be resilient to real-world linguistic variation as it performed well across informal or emotionally expressive language.

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.

H. Production Deployment Considerations

The model strikes a balance between superior performance and enterprise deployment requirements. With an inference latency of 12 ms and a small footprint of 89 MB, it can process in real time without needing GPU infrastructure. It can scale 2,500 complaints/second on commodity servers with support for containerized deployment using Docker and Kubernetes. Its stateless architecture enables adaptive load balancing and secure use — no customer data is stored after processing. Compatibility with REST APIs allows seamless integration with systems such as Salesforce and Zendesk. Model versioning, A/B testing, and automated quality monitoring (latency, accuracy, drift detection) support strong, scalable deployment.

I. Error Analysis and Model Limitations

Most mistakes were between closely related categories (e.g., Data Speed and Network Problems), much due to overlapping vocabulary. Confidence-based routing is suggested: low-confidence predictions (avg. 0.62) can be marked for human evaluation, while high-confidence ones (avg. 0.89) go through automatically. Misclassifications were also found in sarcastic comments, regional spoken dialects, and extremely short complaint texts. Accuracy increased with longer input descriptions, which suggests the importance of richer user input. Shortcomings are the model's existing monolingual configuration and vulnerability to temporal drift as complaint language changes. Retraining schedules and mechanisms for detecting drift will be essential in keeping long-term performance up.

VI. Conclusion and Future Work

Future research will generalize this framework to enable multilingual complaint classification via cross-lingual transfer learning and language-specific attention mechanisms, allowing telecom providers to better serve global user bases. Incorporating speech-to-text systems can further improve real-time call center automation, although this creates challenges such as noise filtering and accent handling. Ensemble transformer models could provide further performance benefits, although system complexity and computational expenses need to be weighed. Temporal modeling of complaint patterns, when correlated with network performance measures, would yield predictive results that could anticipate and anticipate and anticipate initiative-taking service management. Integrating sentiment analysis, urgency prediction, and solution recommender systems will result in fully integrated, intelligent customer service platforms. Multi-task learning would further enhance collective learning across similar customer support tasks. Constructing robust models will be important for field deployment. Model performance can be improved through ensemble techniques of CNN and transformer outputs, or through incorporating structured account metadata using multi-input models. Transfer learning from domain-specific corpora and multilingual training are also promising. This research illustrates the efficiency of transformer-based NLP models—specifically BERT—in telecom complaint classification, registering 92.5% accuracy with enterprise-class feasibility. This method surpasses the conventional models by more than 5%, which corresponds to less manual intervention, enhanced customer satisfaction, and greater resource utilization. Attention-based interpretability also enables human monitoring. This research provides a robust platform for future NLP-based innovations in telecom customer service automation.

References

[1] International Telecommunication Union, "ICT Facts and Figures 2023," Geneva, Switzerland, 2023. [Online]. Available: https://www.itu.int/en/ITU-D/Statistics/Documents/facts/FactsFigures2023.pdf

[2] S. G. Gavrila, C. B. González-Tejero, J. A. G. Gandía, and A. de Lucas Ancillo, "The impact of automation and optimization on customer experience: a consumer perspective," Humanities and Social Sciences Communications, vol. 10, no. 1, pp. 1–10, Nov. 2023. [Online]. Available: https://doi.org/10.1057/s41599-023-02389-0. Accessed: Jun. 7, 2025.

[3] T. Joachims, "Text categorization with Support Vector Machines: Learning with many relevant features," in Machine Learning: ECML-98, C. Nédellec and C. Rouveirol, Eds. Berlin, Heidelberg: Springer, 1998, vol. 1398, pp. 137–142. [Online]. Available: https://doi.org/10.1007/BFb0026683. Accessed: Jun. 7, 2025.

[4] V. Vinayak and J. Chandrasekharan, "Consumer Complaints Classification using Deep Learning & Word Embedding Models," in Proc. 2023 14th Int. Conf. on Computing Communication and Networking Technologies (ICCCNT), 2023, pp. 1–6. [Online]. Available: https://doi.org/10.1109/ICCCNT56998.2023.10307286. Accessed: Jun. 7, 2025

[5] Y. Kim, "Convolutional Neural Networks for Sentence Classification," arXiv preprint arXiv:1408.5882, 2014. [Online]. Available: https://doi.org/10.48550/arXiv.1408.5882. Accessed: Jun. 7, 2025.

[6] Z. Yang, D. Yang, C. Dyer, X. He, A. Smola, and E. Hovy, "Hierarchical Attention Networks for Document Classification," in Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, San Diego, CA, USA, Jun. 2016, pp. 1480–1489. [Online]. Available: https://doi.org/10.18653/v1/N16-1174. Accessed: Jun. 7, 2025.

[7] A. Vaswani et al., "Attention Is All You Need," arXiv preprint arXiv:1706.03762, 2017. [Online]. Available: https://doi.org/10.48550/arXiv.1706.03762. Accessed: Jun. 7, 2025.

[8] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," arXiv preprint arXiv:1810.04805, 2018. [Online]. Available: https://doi.org/10.48550/arXiv.1810.04805. Accessed: Jun. 7, 2025.

[9] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, "RoBERTa: A Robustly Optimized BERT Pretraining Approach," arXiv preprint arXiv:1907.11692, 2019. [Online]. Available: https://doi.org/10.48550/arXiv.1907.11692. Accessed: Jun. 7, 2025.

[10] Y. Lin, H. Chen, W. Xia, F. Lin, Z. Wang, and Y. Liu, "A Comprehensive Survey on Deep Learning Techniques in Educational Data Mining," arXiv preprint arXiv:2309.04761v3, Feb. 2024. [Online]. Available: https://doi.org/10.48550/arXiv.2309.04761. Accessed: Jun. 7, 2025.

[11] S. Gudavalli, "Machine Learning Applications in Telecommunications," SSRN, Sep. 20, 2024. [Online]. Available: https://ssrn.com/abstract=5068351. Accessed: Jun. 7, 2025.

[12] M. Mashaabi, A. Alotaibi, H. Qudaih, R. Alnashwan, and H. Al-Khalifa, "Natural Language Processing in Customer Service: A Systematic Review," arXiv preprint arXiv:2212.09523, Dec. 2022. [Online]. Available: https://doi.org/10.48550/arXiv.2212.09523. Accessed: Jun. 7, 2025.

[13] S. Kingston, Prassanth, S. A. V., B. M. S., and M. K. Rajagopal, "Comparative study on Judgment Text Classification for Transformer Based Models," arXiv preprint arXiv:2306.01739, Apr. 2023. [Online]. Available: https://doi.org/10.48550/arXiv.2306.01739. Accessed: Jun. 7, 2025.

[14] T. Rajapakshe, R. Rana, S. Khalifa, and B. W. Schuller, "Domain Adapting Deep Reinforcement Learning for Real-world Speech Emotion Recognition," arXiv preprint arXiv:2207.12248v3, Dec. 2024. [Online]. Available: https://doi.org/10.48550/arXiv.2207.12248. Accessed: Jun. 7, 2025.

[15] Q. Li, H. Peng, J. Li, C. Xia, R. Yang, L. Sun, P. S. Yu, and L. He, "A Survey on Text Classification: From Traditional to Deep Learning," ACM Transactions on Intelligent Systems and Technology, vol. 13, no. 2, pp. 1–41, Apr. 2022. [Online]. Available: https://doi.org/10.1145/3495162 Accessed: Jun. 7, 2025

Authorized licensed use limited to: National Library of Technology. Downloaded on December 27,2025 at 14:00:35 UTC from IEEE Xplore. Restrictions apply.